

Planning and Reinforcement Learning in AI Systems

(PARLAI)

Lesson 1a - Introduction

Sarah Keren

The Taub Faculty of Computer Science
Technion - Israel Institute of Technology

Course structure and objectives

- I am a member of the Computer Science Faculty at the Technion.
- I head the Collaborative AI and Robotics (CLAIR) lab
- My research focuses on multi-agent AI and multi-robot systems.
- My work offers novel approaches to **increasing the capabilities of AI agents via the design of the environments in which they are intended to act.**
- My email: *sarahk@cs.technion.ac.il*


- **Teaching:** Sarah Keren
- **TA:** Guy Azran
- **Schedule:**
 - Lesson: Wednesday 14:30-16:30
 - Tutorial: Wednesday 16:30-17:30
 - Reception Hours: per request.

What is this class about?

Course structure
and objectives

Decision Making
in AI

Sequential Decision Making: deciding how to act now to accomplish a long-term objective



- From observations to decisions
- From low-level to high-level reasoning
- From theory to practice to theory ...
- From single agent to multi-Agent

There is a wide spectrum of AI research questions.

We focus on:

- Non-adversarial settings.
- Single-agent settings.
- Complex and long-term tasks.
- Combining approaches and principled ways to choose the best approach.

- Weekly 2-hour class followed by a 1-hour tutorial.
- In class, we will cover the theory behind decision-making in AI.
- In the lab, we will practice and test the ideas we discuss in class.
- Grade is set by:
 - 3 homework assignments
 - Quizz (magen)
 - Final exam

All materials will be available online - participation is non-mandatory.

Decision Making in AI

Reactive vs. Rational Agent




Course structure
and objectives

Decision Making
in AI

Reactive vs. Rational Agent


Course structure
and objectives

Decision Making
in AI



- **Reactive or Reflex agent:**
 - simple reactions to stimulus (short-term)
 - fast response times and limited resource usage
- **Rational agent:**
 - extended reasoning capabilities
 - can achieve complex objectives
 - utility maximizing
 - can learn new knowledge and react to change

What does it mean to be autonomous?




We seek a **policy** from states (or knowledge about the state) to actions.

Components of an Autonomous Agent?

Course structure
and objectives

Decision Making
in AI

Action - Perception - Control



Sequential Decision Making

Course structure
and objectives

Decision Making
in AI



VS.



We will focus on sequential decision making (a.k.a. planning)

Sequential Decision Making

Course structure
and objectives

Decision Making
in AI



We will focus on sequential decision making (a.k.a. planning)

Approaches to Decision Making

We seek a **policy**, which we mark as π that maps states $\mathcal{S}$ or (knowledge about the state) to actions $\mathcal{A}$.

Deterministic $\pi : \mathcal{S} \mapsto \mathcal{A}$ or probabilistic $\pi : \mathcal{S} \times \mathcal{A} \mapsto [0, 1]$

More accurately, we seek a mapping from beliefs $\mathcal{B}$, representing knowledge about the current state (formally defined later on), to actions.

Deterministic $\pi : \mathcal{B} \mapsto \mathcal{A}$ or probabilistic $\pi : \mathcal{B} \times \mathcal{A} \mapsto [0, 1]$

Basic approaches:

Approaches to Decision Making

Course structure
and objectives

Decision Making
in AI

We seek a **policy**, which we mark as π that maps states $\mathcal{S}$ or (knowledge about the state) to actions $\mathcal{A}$.

Deterministic $\pi : \mathcal{S} \mapsto \mathcal{A}$ or probabilistic $\pi : \mathcal{S} \times \mathcal{A} \mapsto [0, 1]$

More accurately, we seek a mapping from beliefs $\mathcal{B}$, representing knowledge about the current state (formally defined later on), to actions.

Deterministic $\pi : \mathcal{B} \mapsto \mathcal{A}$ or probabilistic $\pi : \mathcal{B} \times \mathcal{A} \mapsto [0, 1]$

Basic approaches:

- Programming/reactive
- Model-Based Planning
- Machine Learning (from data and experience)
- Reinforcement Learning

Example

You are in Kibutz Ein Dor and you want to get to Austin Texas.



How do you decide how to get to Austin?

Course structure
and objectives

Decision Making
in AI




Very practical in many applications, but lacks flexibility.

Model-Based Planning

Course structure
and objectives

Decision Making
in AI



Key idea: use the same language for all planning domains.

How can we create the same language for logistics, Mars rover exploration, and cybersecurity?

Characterized by:

- **States** $\mathcal{S}$ a system can be in.
 - A state is a full assignment to the set of **variables** (features) $\mathcal{X}$.
- **Actions** $\mathcal{A}$ that change the current state by changing the values of certain variables.
- **Reward Function** $\mathcal{R}$ that sets a numeric signal passed from the environment (can represent cost)
 - used to signal the objective
 - some domains have **goal conditions** such that the agent should reach **goal states** (e.g., 'be at Austin').
- **Objective:** find a **policy** that drives the initial state to a goal state or that maximizes the expected accumulated reward.

Use **statistics** to find patterns in massive amounts of data.

The input vector x (e.g., an image) is mapped to an output $f(x)$, which represents a classification label or a probability distribution over the possible labels. Paradigms of Machine Learning

- **Supervised Learning**

- Labeled data (x, y)
- Classification, regression

- **Unsupervised Learning**


- No labels
- Clustering, representation learning

- **Reinforcement Learning**

- Interaction with environment
- Sequential decisions, delayed rewards

Template algorithm (for supervised learning):

- **Input:**
 - Sample set: $S = \{(x_i, y_i)_{i=1}^m\}$
 - Model calls H (candidate classifiers)
- **Output:**
 - Classifier with lowest **expected error** (e.g. loss).



Machine Learning: Imitation Learning

Course structure
and objectives

Decision Making
in AI

- Learning from demonstrations
- Useful when it is easier for an expert to demonstrate the desired behavior rather than to specify a reward function.
- Simplest form: **behaviour cloning (BC)**, learning the expert's policy using supervised learning.




What's the sample set here? What does error (loss) represent?

What Changes with Foundation Models?

In the (not so long ago) past

- Train **one model per task**
- Requires labeled data or reward design
- Limited generalization
- Start from scratch each time

With foundation models, we move from task-specific learning to general-purpose models.




Foundation Models: A New Learning Paradigm

Course structure
and objectives

Decision Making
in AI

- **Core idea:** train once, reuse everywhere
- **Pretraining**
 - Massive, diverse data
 - Self-supervised objectives
 - Learn general representations
- **Adaptation**
 - Fine-tuning, prompting, adapters, RL
 - Task- and domain-specific objectives



Are We Done?

Course structure
and objectives

Decision Making
in AI




Even if we have advanced planning, classification, and imitation learning methods - what's missing?

Reinforcement Learning (RL)

Course structure
and objectives

Decision Making
in AI

- **Learning from interaction** - how to behave to maximize the accumulated reward/achieve a goal.
- The learner is not told which actions to take, but instead **must discover** which actions yield the highest total **reward** by trying them.



Is reward enough?

<https://deepmind.com/research/publications/2021/Reward-is-Enough>

Reinforcement Learning (contd.)


Course structure
and objectives

Decision Making
in AI

The process of Reinforcement Learning (RL) involves:

- Observing the environment
- Deciding how to act using some strategy (policy)
- Acting accordingly
- Receiving a reward (or penalty) and an observation
- Learning from the experiences and refining the policy

The policy is iteratively improved until convergence to an optimal policy.



- RL vs. Planning?
- RL vs. ML ?
- RL vs. IL ?

RL vs. Model-Based Planning

Course structure
and objectives


Decision Making
in AI

Both are fundamental problems in sequential decision-making that aim to find a sequence of actions to achieve a specific goal or maximize utility

- **Model-Based Planning:** based on a known model of the environment and its dynamics.
- **Reinforcement Learning:** model of the environment is initially not fully known. The agent improves its policy by interacting with the environment.

The distinction is very fuzzy...

- There is no supervisor, only a reward signal
- Feedback is delayed, not instantaneous
- Data is sequential (non i.i.d): agent's actions affect the subsequent rewards it receives



Both address sequential decision-making.

- **Imitation Learning (IL):**

Policy learned directly from expert demonstrations, treating the problem as supervised (or structured prediction) learning, without explicit reward optimization.

- **Reinforcement Learning (RL):**


Policy learned by interacting with the environment and optimizing a reward signal.

Summary of AI approaches

Course structure
and objectives

Decision Making
in AI

- Different approaches for decision-making: from data-driven to model-based approaches.
- Different settings require different solutions.



Open Questions:

- How to choose the right paradigm for your problem?
- What is a good model for the decision making process?
- How does the presence of **other agents** affect these decision?